

1 CHAPTER 01 : FOUNDATIONS OF DEEP LEARNING

1.1 MCCULLOCH-PITTS NEURON




Figure 1: McCulloch-Pitts Neuron

Equations :

$$\begin{cases} y = g(f(x_1, x_2, \dots, x_n)) \\ f(x_1, x_2, \dots, x_n) = \sum_{i=1}^n w_i * x_i \\ g(z) = \begin{cases} 1 & \text{if } z \geq \theta \\ 0 & \text{otherwise} \end{cases} \end{cases}$$

Characteristics :

- weights and threshold are set manually (no learning).
- can not solve non-linear problems.
- takes only boolean inputs.

1.2 PERCEPTRON




Figure 2: Perceptron

Equations :

$$\begin{cases} y = g(f(x_1, x_2, \dots, x_n)) \\ f(x_1, x_2, \dots, x_n) = W^t * x + b = \sum_{i=1}^n w_i * x_i + b \\ g(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases} \end{cases}$$

g is called the *Heaviside* activation function.

Characteristics :

- weights and threshold are learned.
- can not solve non-linear problems.
- handles only classification problems.

Algorithm 1 Perceptron's learning algorithm

```

for  $i \leftarrow 1$  to  $L$  do
  for  $j \leftarrow 1$  to  $m$  do
     $error \leftarrow y_j - g(w^T * x_j + b)$ 
     $w \leftarrow w + r \times error \times x_j$ 
     $b \leftarrow b + r \times error$ 
  end for
end for

```


1.3 ADALINE

Figure 3: Adaline

the main difference between **Adaline** and **Perceptron** is in the learning process, adaline uses the continuous values (before applying Heaviside/Quantizer) to update the weights.

Equations :

$$\begin{cases}
 y = g(f(x_1, x_2, \dots, x_n)) \\
 f(x_1, x_2, \dots, x_n) = W^t * x + b = \sum_{i=1}^n w_i * x_i + b \\
 g(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}
 \end{cases}$$

Characteristics :

- weights and threshold are learned.
- can not solve non-linear problems.
- multi-layer perception is equivalent to a simple linear regression.

Algorithm 2 Adaline's learning algorithm

```

for  $i \leftarrow 1$  to  $L$  do
  for  $j \leftarrow 1$  to  $m$  do
     $error \leftarrow y_j - f(x_j)$ 
     $w \leftarrow w + \alpha \times error \times x_j$ 
     $b \leftarrow b + \alpha \times error$ 
  end for
end for

```

1.4 MULTI-LAYER PERCEPTRON

1.4.1 HOW CAN A MULTI-LAYER PERCEPTRON HELP SOLVE MORE COMPLEX TASKS ? : XOR EXAMPLE




Figure 4: Multi-layer perceptron

x_1	x_2	n_1		n_2		o_1	
		z	y	z	y	z	y
0	0	-0.5	0	-1.5	0	-0.5	0
1	0	0.5	1	-0.5	0	0.5	1
0	1	0.5	1	-0.5	0	0.5	1
1	1	1.5	1	0.5	1	-0.5	0

1.4.2 NEURAL NETWORKS CHARACTERISTICS

- **size** : number of the nodes in the model.
- **width** : number of nodes in a specific layer.
- **depth** : number of layers in the neural network.
- **architecture** : the specific arrangement of the layers and nodes in the model.

1.4.3 HOW TO DESIGN A NEURAL NETWORK ARCHITECTURE ?

- **Experimentation.**
- **Intuition** : comes from experience in a specific domain.
- **Go for depth** : deeper neural networks can help solve complex problems but are more vulnerable to overfitting.
- **Borrow ideas** : architectures that are proved to work well on similar problems is a good starting point.
- **Search** : Heuristics, Random search.


1.4.4 NEURAL NETWORKS : FORWARD PASS

Notations :

- L : the number of layers.
- g_l : the activation function of the l^{th} layer.
- $a^{[l]}$: the output of the l^{th} layer.
- $z^{[l]}$: the output of the l^{th} layer before applying the activation function.
- $a_n^{[l]}$: the output of the n^{th} neuron of the l^{th} layer.
- θ_l : the weights of the l^{th} layer.

$$X = \begin{bmatrix} x_0 & x_1 & x_2 \\ 1 & 0.504 & -0.4161 \\ 1 & -0.99 & -0.6536 \\ 1 & 0.2837 & 0.9602 \end{bmatrix}$$

x_0 is always equal to 1 because it represents the bias.



FORWARD EQUATIONS

$$\begin{cases} a^{[1]} = X^t \\ a^{[l+1]} = g_l(\theta_{l+1} \times a^{[l]}) + \text{add a row of ones} \\ \hat{y} = (a^L)^t + \text{without the additional row of ones} \end{cases}$$

LOSS

Cross entropy loss :

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \sum_{j=1}^K [-y_j^{(i)} * \log(\hat{y}_j^{(i)})]$$

Binary cross entropy loss :

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} * \log(\hat{y}^{(i)}) - (1 - y^{(i)}) * \log(1 - \hat{y}^{(i)})]$$

BACKWARD EQUATIONS FOR AN MLP WITH ONE HIDDEN LAYER

$$\begin{cases} \frac{\partial J}{\partial \theta_2} = \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z_2} \frac{\partial z_2}{\partial \theta_2} \\ \frac{\partial J}{\partial z_2} = \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z_2} = \hat{y} - y \text{ if the activation of the last layer is sigmoid or softmax} \\ \frac{\partial z_2}{\partial \theta_2} = a_1 \text{ the output of hidden layer} \end{cases}$$

$$\begin{cases} \frac{\partial J}{\partial \theta_1} = \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z_2} \frac{\partial z_2}{\partial a_1} \frac{\partial a_1}{\partial z_1} \frac{\partial z_1}{\partial \theta_1} \\ \frac{\partial J}{\partial z_2} = \frac{\partial J}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial z_2} = \hat{y} - y \text{ if the activation of the last layer is sigmoid or softmax} \\ \frac{\partial z_2}{\partial a_1} = \theta_2 \\ \frac{\partial a_1}{\partial z_1} = \frac{\partial g_1}{\partial t}(t) \text{ depends on the hidden's layer activation function} \\ \frac{\partial z_1}{\partial \theta_1} = x \text{ the input data} \end{cases}$$

1.4.5 ACTIVATION FUNCTIONS

WHY USE NON-LINEAR ACTIVATION FUNCTIONS ?

Because a neural network with a linear activation function is equivalent to a simple linear regression model, so it will perform poorly on non-linear separable problems, linear activation is only used in the output layer in regression problems.

COMMON ACTIVATION FUNCTIONS

name	expression	derivative	notes
Sigmoid	$\sigma(x) = \frac{1}{1+e^{-x}}$	$\sigma(x) * (1 - \sigma(x))$	- used in hidden and output layer. - can cause vanishing gradient descent problem. - expensive to compute
ReLU	$relu(x) = \max(0, x)$	$\begin{cases} 1 & \text{if } x \geq 0 \\ 0 & \text{otherwise.} \end{cases}$	- used in hidden layers. - can cause dying relu problem. - doesn't cause vanishing gradient descent. - easy to compute
Tanh	$\tanh(x) = \frac{e^x + e^{-x}}{e^x - e^{-x}}$	$1 - \tanh^2(x)$	
Softmax	$softmax(x) = \frac{e^x}{\sum_{i=1}^n e^x}$		used in the output layer of multi-classifications problems.

Table 1: Common activation functions

2 CHAPTER 02 : OPTIMIZING DEEP NEURAL NETWORKS

2.1 BATCH GRADIENT DESCENT VS STOCHASTIC GRADIENT DESCENT VS MINI-BATCH GRADIENT DESCENT

Batch GD	Stochastic GD	Mini-Batch GD
Processes all the dataset each iteration	processes one sample each iteration	process a portion (batch) of the dataset each iteration
$batch_size = m$	$batch_size = 1$	$batch_size$ is a hyper-parameter.

	Advantages	Disadvantages
Batch GD	<ul style="list-style-type: none"> - guaranteed to converge in theory - unbiased estimate of the gradient 	<ul style="list-style-type: none"> - slow for large datasets - memory issues for large datasets
Mini-Batch GD	<ul style="list-style-type: none"> - Faster than Batch GD - adds noise which can help improve generalization 	<ul style="list-style-type: none"> - can cause oscillations - may require learning rate decay
Stochastic GD	<ul style="list-style-type: none"> - Same as Mini-Batch 	<ul style="list-style-type: none"> - slow run time - adds a lot of noise

Classic SGD can cause oscillation and prevent using larger learning rates making convergence slower.

2.2 SGD WITH MOMENTUM

$$\begin{cases} W_{t+1} = W_t - \alpha * V_t \\ V_t = \beta * V_t - (1 - \beta) * \Delta W_t \end{cases}$$

the default value for β is **0.9**.

2.3 RMSPROP OPTIMIZER

$$\begin{cases} W_t = W_{t-1} - \alpha * \frac{\Delta w_t}{\sqrt{V_t + \epsilon}} \\ V_t = \beta * V_{t-1} + (1 - \beta) * \Delta W_t^2 \end{cases}$$

the default values for β and ϵ are **0.999** and 10^{-8} respectively.

2.4 ADAM OPTIMIZER

$$\begin{cases} W_t = W_{t-1} - \alpha * \frac{V_t}{\sqrt{S_t + \epsilon}} \\ S_t = \beta_2 * S_{t-1} + (1 - \beta_2) * \Delta W_t^2 \\ V_t = \beta_1 * V_{t-1} + (1 - \beta_1) * \Delta W_t \end{cases}$$

the default values for β_1, β_2 and ϵ are 0.9, 0.999 and 10^{-8} respectively.

2.5 LEARNING RATE DECAY

Time based decay :

$$\alpha_t = \alpha_0 \frac{1}{1 + decay_rate * t}$$

Step based decay :


$$\alpha_t = \alpha_0 * drop_rate^{t/epoch_drop}$$

Exponential decay :

$$\alpha_t = \alpha_0 * e^{-decay_rate * t}$$

2.6 OVERFITTING & REGULARIZATION TECHNIQUES

2.6.1 OVERFITTING VS MODEL COMPLEXITY



2.6.2 L1 & L2 NORMALIZATION

L2 Norm (weight decay)

$$\begin{cases} J(\theta) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}, y) + \frac{\lambda}{2 \cdot m} \sum_{l=1}^L (\|W^l\|)^2 = E(\theta) + \frac{\lambda}{2 \cdot m} \sum_{l=1}^L \|W^l\|^2 \\ \|W^l\|^2 = \sum_{i,j} w_{i,j}^{(l)2} \\ \frac{\partial J}{\partial W} = \frac{\partial E}{\partial W} + \frac{\lambda}{m} * W \end{cases}$$

L1 Norm

$$\begin{cases} J(\theta) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}, y) + \frac{\lambda}{2 \cdot m} \sum_{l=1}^L (\|W^l\|) = E(\theta) + \frac{\lambda}{2 \cdot m} \sum_{l=1}^L (\|W^l\|) \\ \|W^l\| = \sum_{i,j} |w_{i,j}| \end{cases}$$

2.6.3 DROPOUT :

Dropout is a regularization technique that randomly turns off some neurons each iteration, the number of the neurons to mask is controlled by the **dropout rate** parameter, the dropout is deactivated during inference.

2.6.4 EARLY STOPPING :

- Monitoring model performance : the choice of metric, validation set.
- Trigger to Stop : stop the training when the loss increases or became unstable for a certain number of epochs.
- The choice of the model : save the weights each time the loss decreases.

2.6.5 BATCH NORMALIZATION :

Batch normalization is a regularization technique that speeds up training and handles internal covariant shift, batch normalization is an extra layer that normalizes the batch by subtracting its mean and dividing it by its standard deviation.

2.7 GRADIENT CHECKING

Gradient checking is a technique used to verify whether the implementation of the backpropagation is true or not by calculating an approximation of the derivatives and comparing it with the results of the backpropagation.

$$\frac{\partial J}{\partial \theta} \approx \frac{J(\theta + \epsilon) - J(\theta - \epsilon)}{2 * \epsilon}$$

2.8 HYPERPARAMETER TUNING

Hyperparameter tuning technique aims to find a better combination of hyperparameters that leads to a better performance and less overfitting, whether those parameters are related to the network's architecture or the optimizer.

- Manual search : babysitting.
- random search.
- grid search.
- Bayesian optimization.

GRID SEARCH VS RANDOM SEARCH

Random Search	Grid Search
Doesn't guarantee to find the best hyperparameters	guarantees to find the best hyperparameters
Pick random points to try from the configuration space	tries all the possible combinations.
Good in high spaces	Curse of dimensionality
Good results in less iterations	computationally expensive.

Table 2: Grid search Vs Random Search

Grid search & Random search share a common downside : "each guess is independent from the previous one", solution : **Bayesian optimization.**

3 CHAPTER 03 : CONVOLUTIONAL NEURAL NETWORKS

3.1 Foundations of CNNs

Convolutional neural networks are a type of deep learning architectures designed at first for computer vision tasks but later proved to work on other fields like text and voice processing.

A convolutional neural network mainly consists of three type of layers :

- Convolutional layers : for feature extraction.
- Pooling layers : reduce the size of the feature map.
- Fully connected layers : decision making.

3.1.1 Convolutional layer

A convolutional layer takes as an input a **feature map** of size : $W \times H \times C$, C is called the number of channels, and applies M different filters of dimensions $F_w \times F_h$ to produce a new feature map of dimension : $W' \times H' \times M$.

The convolutional layer processes the input feature map using a moving window of size $F_w \times F_h$, the window moves by S steps, S is called the stride.

Optionally, before passing the feature map to the convolutional layer a *padding* can be added by adding additional layers of pixels around the border of an image it can serve two main purposes : **dimension preservation (same padding)** and **preventing information loss around the borders**.




Figure 5: Convolutional layer in action

Calculating the size of the new feature map :

For simplification : $F_w = F_h = F$ and $W = H = N$ and $W' = H' = N'$.

$$N' = \frac{N - F + 2 * P}{S} + 1$$

Calculating the number of parameters :

$$F^2 \times C \times M + M$$

Example :

$$P = 1, S = 2, F = 3, C = C_{in} = 1, M = C_{out} = 1, N = 3$$

The size of the new feature map :

$$N' = \frac{3 - 3 + 2 * 1}{2} + 1 = 2$$

Numerical demonstration :

$$X = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 125 & 15 & 30 & 0 \\ 0 & 250 & 25 & 2 & 0 \\ 0 & 174 & 255 & 10 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$F = \begin{bmatrix} 0 & 0 & 1 \\ 0 & -1 & 0 \\ 1 & 0 & 0 \end{bmatrix}$$

$$Y = X \odot F = \begin{bmatrix} -125 & -5 \\ -149 & -10 \end{bmatrix}$$

3.1.2 Pooling layer

Pooling layers help in reducing the dimensions of the feature maps while retaining the most important features.

- Max pooling : takes the maximum of the region.
- Average pooling : takes the average of the region.
- Global max pooling: takes the maximum value over the entire feature map.
- Global average pooling: takes the average value over the entire feature map.
- L2 pooling: takes the L2 norm of each region.
- Fractional max pooling: takes the maximum value over a randomly chosen subset of the region.

The dimensions of the resulted feature map is calculated the same way as the convolutional layer, because pooling layers also process the input feature map using a moving window.




Figure 6: Max-Pooling layer in action

Example :

$$P = 1, S = 1, F = 2$$

$$X = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 125 & 15 & 30 & 0 \\ 0 & 250 & 25 & 2 & 0 \\ 0 & 174 & 255 & 10 & 0 \\ 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

$$Max - Pool(X) = \begin{bmatrix} 125 & 125 & 30 & 30 \\ 250 & 250 & 30 & 30 \\ 250 & 255 & 255 & 10 \\ 174 & 255 & 255 & 10 \end{bmatrix}$$

3.2 CNN vs ML

	CNN	ML
Data Requirement	Large amount of labeled data	small amount of data
Feature extraction	Automatic	Manual or using unsupervised algorithms
Training	requires computational resources	doesn't requires computational resources
Generalization	can generalize well on unseen data	may struggle on new data

Table 3: CNN vs ML

3.3 Deep convoloutions architectures

3.3.1 LeNet

A classic CNN architecture composed of 5 convolution layers and 3 fully connected layers the size of the filter is fixed to 5.

3.3.2 AlexNet

The AlexNet architecture consists of 5 convolutional layers, followed by 3 fully connected layers, despite its small depth it has more than 56 millions parameters.

3.3.3 ZFNET

ZFNET uses smaller filters sizes which helps in both capturing more fine-grained features and reducing the number of parameters.

3.3.4 VGGNet

VGGNet architecture uses stacked convolution layers of smaller filters sizes (3×3), because it reduces the number of parameters allowing the neural network to be deeper and adds more non-linearity to the architecture.

3.3.5 Inception (GoogleNet)

The Inception net architecture passes the input through multiple convolutional layers with various filter sizes to extract both low-level and high-level features. It then concatenates these features, making this architecture suitable for feature extraction, and to overcome the potential large number of parameters it also add 1×1 conv-layers before the convolution and after the max-pool layer.




Figure 7: Inception Layer

3.3.6 ResNet

Other than the large number of parameters another challenge that arises when training very deep neural networks is the problem of Vanishing/Exploding gradient.

To overcome this problem, residual blocks were introduced, the idea is that even if $F(x) \approx 0$, then the input x is still propagated through the network allowing the last layers to also learn.




Figure 8: Residual Block

3.3.7 MobileNet

To reduce the number of parameters the mobile-net architecture divides a simple convolution layer into a depth and a point wise convolutions. the depth-wise convolution are used to apply a **single** filter to into each input channel. while the point-wise convolution applies a 1×1 **conv-layer** to the output of the Depth-wise convolution to create new features.

Standard convolution cost :

$$C_{in} \times C_{out} \times N \times N \times F \times F$$

Depth-wise convolution cost :

$$C_{in} \times N \times N \times F \times F$$

Point-wise convolution cost :

$$C_{in} \times C_{out} \times N \times N$$

Reduction in cost :

$$\frac{C_{in} \times N \times N \times F \times F + C_{in} \times C_{out} \times N \times N}{C_{in} \times C_{out} \times N \times N \times F \times F} = \frac{1}{C_{out}} + \frac{1}{F^2}$$

3.3.8 R-CNN

R-CNN is a deep neural network used for object detection, it proposes a set of regions using **selective search** then a CNN model to classify these regions.

3.4 Transfer learning

In machine learning, transfer learning is the reuse of a pre-trained model on a new problem.

3.4.1 Approaches

- Training the model to reuse it.
- Using a pretrained model.

3.4.2 How can pre-trained models be used ?

- As feature extractors.
- Fine tuning a subset of layers.
- Fine tuning all layers.

4 CHAPTER 04 : DIMENSIONALITY REDUCTION WITH DEEP LEARNING

4.1 Auto-encoders and unsupervised learning

Auto-encoders are a type neural networks that are used to learn data-encoding,by learning an approximation of the identity function,an auto-encoder consists of three parts :

- Encoder : maps the higher dimensional input to a lower dimensional representation.
- Bottleneck : contains the lower dimensional representation.
- Decoder : tries to reconstruct the original image from lower dimensional representation.

4.1.1 Types of auto-encoders

To train auto-encoders usually MSE and L1 loss are used.

Undercomplete autoencoders :

Undercomplete autoencoders tries to predict the exact same input,they are used for dimentionality reduction and usually performs better than PCA because it can learn non-linear relationships.

Sparse Autoencoders :

sparse auto-encoders impose **sparsity-constraint** to prevent overfitting.

$$\begin{cases} L_{new} = L + \beta * \sum_{j=1}^s KL(\rho \parallel \hat{\rho}_j) \\ \hat{\rho}_j = \frac{1}{m} \sum_{i=1}^m a_j x^{(i)} \\ KL(\rho \parallel \hat{\rho}_j) = \rho \log(\frac{\rho}{\hat{\rho}_j}) + (1 - \rho) \log(1 - \frac{\rho}{1 - \hat{\rho}_j}) \end{cases}$$

ρ is the sparsity parameter and β is the regularization parameter.

Denoising Autoencoders :

Denoising autoencoders are trained to remove noise from data,it takes as input a noisy input (introduced for example by randomly swapping some pixels) and the network tries to predict the real (without noise) image.

Stacked autoencoders :

Stacked autoencoders consists of multiple auto-encoders stacked on top of each other,used for complex datasets.

4.1.2 Auto-encoders applications

- Dimensionality reduction.
- Noise reduction.

4.2 Generative adversarial networks (GANs)

GANs are a type of neural network architecture that are used in content generation (images, text...etc), it consists of two main blocks (networks), the **generator** and the **discriminator**.

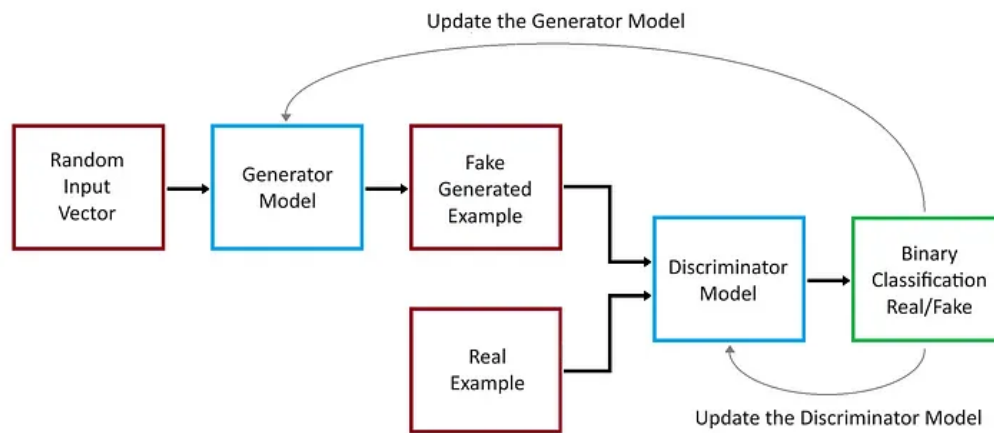


Figure 9: GANs

Generator : This network takes random noise as input and produces data (like images). Its goal is to generate data that's as close as possible to real data.

Discriminator : This network takes real data and the data generated by the Generator as input and attempts to distinguish between the two. It outputs the probability that the given data is real.

4.2.1 Training process

1. Define the GAN architecture.
2. Train discriminator on real dataset.
3. Generate fake inputs from the generator.
4. Train discriminator on fake data.
5. Generate new fake data.
6. Repeat 2,3,4,5 steps for multiple epochs.

4.2.2 Types of GANs

- Vanilla GAN.
- Deep Convolutional GAN (DCGAN).
- Conditional GAN (CGAN).
- Laplacian Pyramid GAN (LAPGAN)